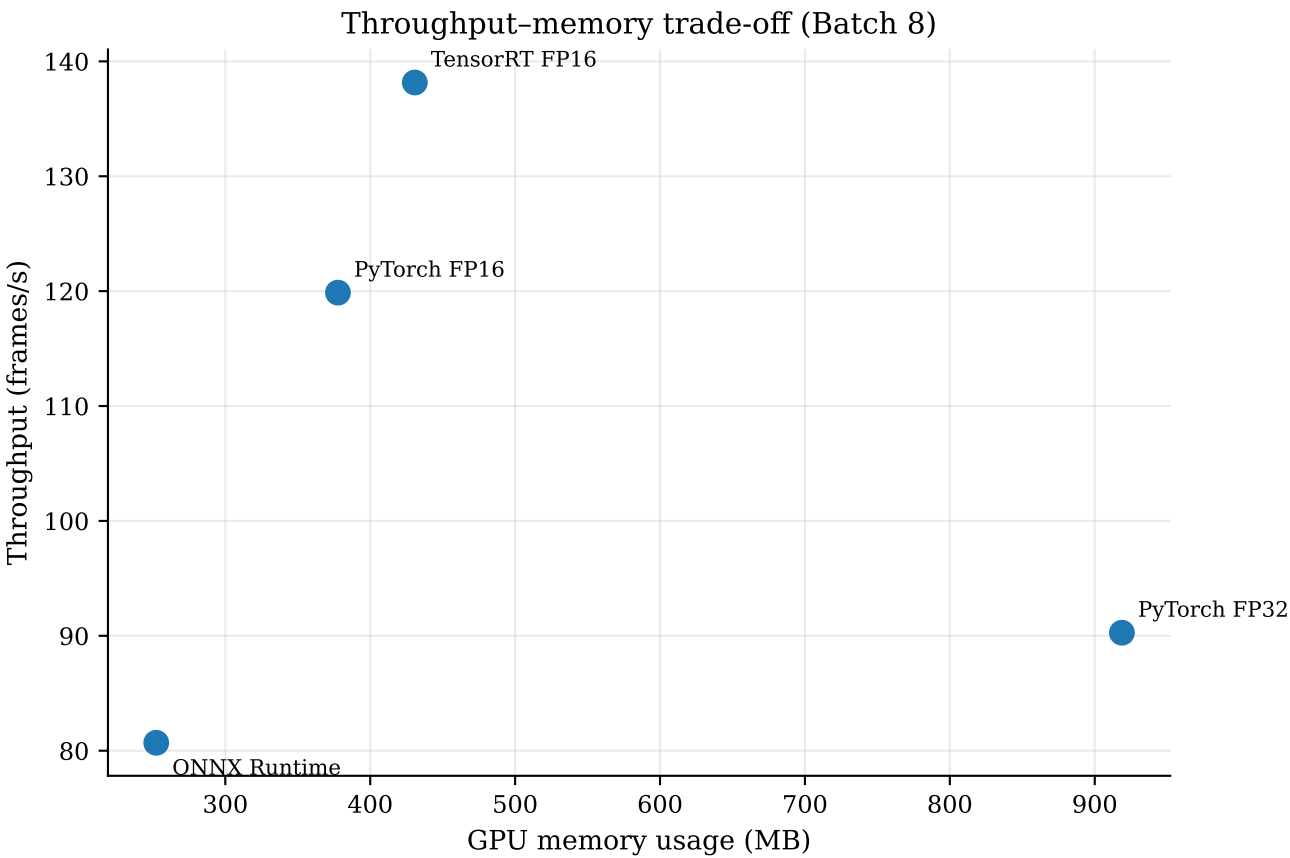
Throughput–memory trade-off (Batch 8)
Throughput (frames/s)
GPU memory usage (MB)
140
130
120
110
100
90
80
300
400
500
600
700
800
900
TensorRT FP16
PyTorch FP16
PyTorch FP32
ONNX Runtime